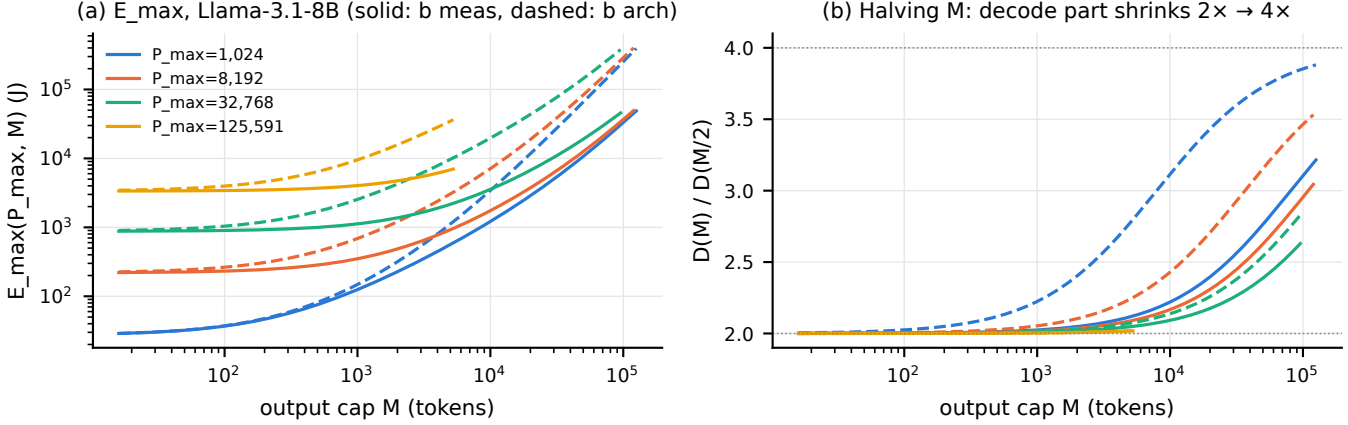
(a) E_max, Llama-3.1-8B (solid: b meas, dashed: b arch)
P_max=1,024
P_max=8,192
P_max=32,768
P_max=125,591
E_max(P_max, M) (J)
output cap M (tokens)
(b) Halving M: decode part shrinks 2× → 4×
D(M) / D(M/2)
output cap M (tokens)